

Lyft Data Contest

Emeka Mbazor

9/14/2019

```
# importing libraries
library(tidyverse)

## Registered S3 methods overwritten by 'ggplot2':
##   method      from
##   [.quosures   rlang
##   c.quosures   rlang
##   print.quosures rlang

## -- Attaching packages ----- tidyverse 1.2.1 --

## v ggplot2 3.1.1    v purrr  0.3.2
## v tibble  2.1.3    v dplyr  0.8.1
## v tidyr   0.8.3    v stringr 1.4.0
## v readr   1.3.1    v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

library(dplyr)
library(ggplot2)
library(broom)
library(scales)

##
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
##
##   discard

## The following object is masked from 'package:readr':
##
##   col_factor

library(lubridate)

##
## Attaching package: 'lubridate'

## The following object is masked from 'package:base':
##
##   date

library(modelr)

##
## Attaching package: 'modelr'

## The following object is masked from 'package:broom':
##
##   bootstrap
```

```

library(reshape2)

##
## Attaching package: 'reshape2'
## The following object is masked from 'package:tidyr':
##
##      smiths
# plot settings
theme_set(theme_bw())
options(repr.plot.width=4, repr.plot.height=3)

# loading in datasets
driver <- read.csv("driver_ids.csv")
ride <- read.csv("ride_ids.csv")
ride_stamps <- read.csv("ride_timestamps.csv")

str(driver$driver_onboard_date)    # 49 factors

## Factor w/ 49 levels "2016-03-28 00:00:00",...: 2 2 9 27 33 40 11 41 30 40 ...
str(ride_stamps$timestamp)        # 865827 factors

## Factor w/ 865827 levels "","2016-03-28 05:48:18",...: 702756 702760 702800 702801 703012 335582 335582 ...
# factors as a data type is not particularly helpful...

# converting driver onboard dates from factors to dates
driver <- driver %>%
  mutate(driver_onboard_date2 = as.Date(as.character(driver_onboard_date)))

# converting ride timestamps from factors to dates
ride_stamps <- ride_stamps %>%
  mutate(timestamp2 = as.Date(as.character(timestamp)))

# converting ride timestamps from factors to date-times
ride_stamps <- ride_stamps %>%
  mutate(timestamp3 = ymd_hms(as.character(timestamp)))

# converting date-times to times,
# hms::as.hms doesn't account for timezones,
# leading to a loss of 4 hours
# 4 hours is added back in after the conversion
four_hours <- 60 * 60 * 4 # four hours in seconds
ride_stamps <- ride_stamps %>%
  mutate(timestamp4 = timestamp3 + four_hours) %>%
  mutate(times = hms::as.hms(timestamp4))

# joining all datasets together
ride_info <- left_join(driver,ride, by = "driver_id")

## Warning: Column `driver_id` joining factors with different levels, coercing
## to character vector
ride_info <- left_join(ride_info,ride_stamps, by = "ride_id")

## Warning: Column `ride_id` joining factors with different levels, coercing

```

```

## to character vector
base_fare <- 2
cost_per_mile <- 1.15
cost_per_min <- 0.22
service_fee <- 1.75

# calculating the revenue gained per trip with the equation:
# ((base_fare + (cost_per_mile * ride_distance_miles) + (cost_per_min * ride_duration_min))) * (1 + prime_time_decimal)
ride_info_enhanced <- ride_info %>%
  mutate(ride_duration_min = ride_duration / 60.0) %>%
  mutate(ride_distance_miles = ride_distance / 1609.34) %>%
  mutate(prime_decimal = ride_prime_time / 100) %>%
  mutate(price_per_trip = ((base_fare + (cost_per_mile * ride_distance_miles) + (cost_per_min * ride_duration_min)) * (1 + prime_decimal)))

# rides are always at least 5 dollars and are
# capped at 400 dollars
ride_info_enhanced <- ride_info_enhanced %>%
  mutate(price_per_trip_adj = case_when(
    (price_per_trip > 400) ~ 400,
    (price_per_trip < 5) ~ 5,
    TRUE ~ price_per_trip
  ))

# price per trip info
summary(ride_info_enhanced$price_per_trip_adj)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.      NA's
##    5.000   8.037  10.568  13.538  15.111 400.000       83

#      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.      NA's
#    5.000   8.037  10.568  13.538  15.111 400.000       83

# grouped by driver and summed up ride prices to get
# each rider's revenue
ride_info_drivers <- ride_info_enhanced %>%
  group_by(driver_id) %>%
  mutate(driver_revenue = sum(price_per_trip_adj)) %>%
  filter(!is.na(driver_revenue)) %>%
  ungroup()

# NOTE: We're calculating revenue generated by the drivers but
#       it should still be directly proportional to Lyft's revenue.

# driver revenue info
summary(ride_info_drivers$driver_revenue)

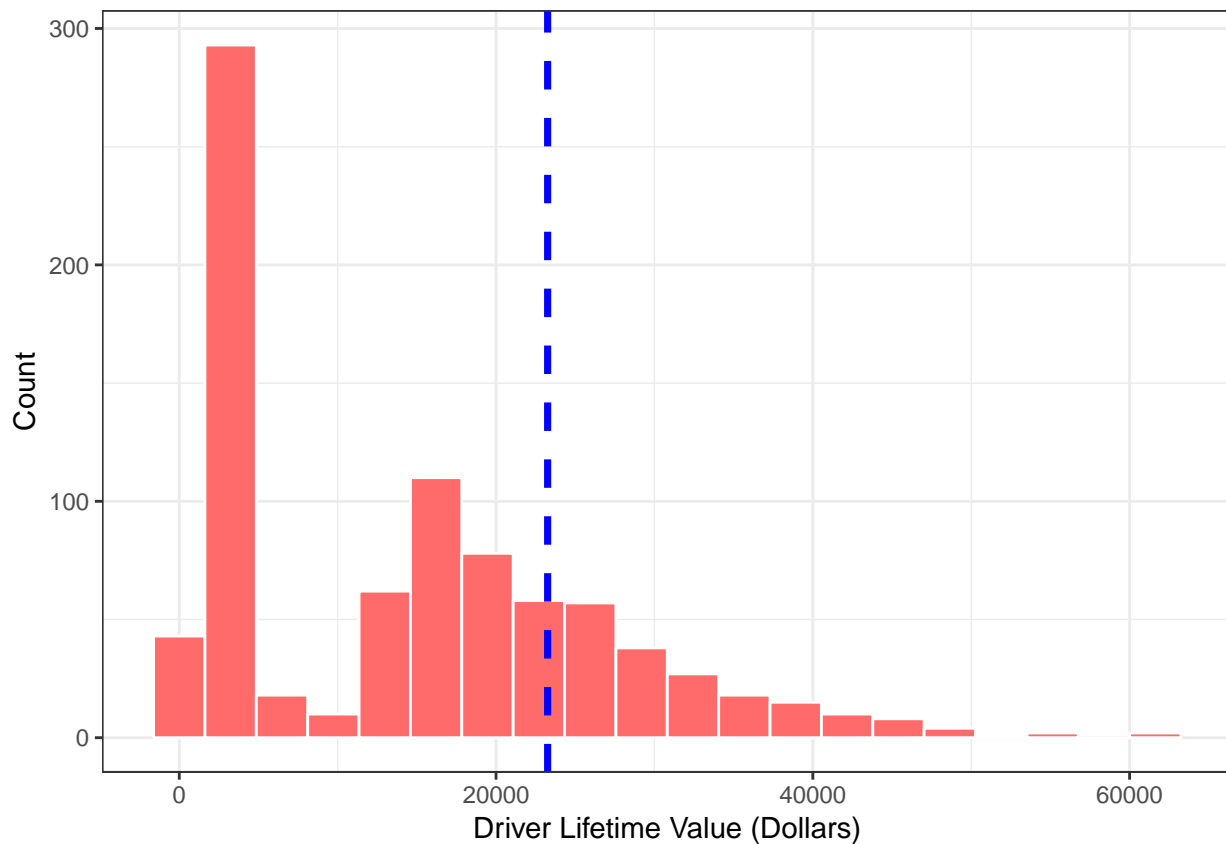
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    128.5 16761.5 23256.6 24502.8 31698.1 61751.4

#      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#    128.5 16761.5 23256.6 24502.8 31698.1 61751.4

# there seems to be a lot of drivers that don't generate much
# revenue at all...
driver_revenue_median = median(ride_info_drivers$driver_revenue)

```

```
ride_info_drivers %>%
  group_by(driver_id) %>%
  summarize(driver_revenue = max(driver_revenue))%>%
  ggplot() +
  geom_histogram(aes(x = driver_revenue),
                 color = "white",
                 fill = "indianred1",
                 bins = 20) +
  geom_vline(xintercept = driver_revenue_median, lwd = 1.25,
             linetype = "dashed", color = "blue") +
  xlab("Driver Lifetime Value (Dollars)") +
  ylab("Count")
```



```
# driver revenue cumulative distribution function

ride_info_drivers <- ride_info_drivers %>%
  group_by(driver_id) %>%
  mutate(num_rides = n())


cdf <- ride_info_drivers %>%
  select(driver_id, driver_revenue, num_rides) %>%
  group_by(driver_id) %>%
  summarize(driver_revenue = max(driver_revenue),
            num_rides = max(num_rides)) %>%
  arrange(desc(num_rides)) %>%
  unique() %>%
```

```

ungroup() %>%
mutate(rank = rank(desc(num_rides)),
       cdf = cumsum(driver_revenue) / sum(driver_revenue)
       )
# a majority of the accumulative driver lifetime value is
# earned by drivers who rank 500 and lower in terms of number of rides
# The decrease in the slope also suggests that past this point
# it's economically inefficient to encourage drivers to
# make more rides

cdf %>%
ggplot(aes(x = rank, y = cdf)) +
  geom_line() +
  scale_y_continuous(label = percent) +
  ylab('Cumulative Percent of Driver') +
  xlab('Driver Rank by Number of Rides With Lyft') +
  ggtitle('Cumulative Driver Lifetime Value')

```



```

### Data Cleaning + Data Prep

## factor generation:

# driver revenue by week
weekly_driver_revenue <- ride_info_enhanced %>%
mutate(event = as.character(event)) %>%
filter(event == "accepted_at") %>%

```

```

mutate(week = week(timestamp2)) %>%
group_by(driver_id, week) %>%
mutate(weekly_driver_revenue = sum(price_per_trip_adj)) %>%
ungroup() %>%
group_by(driver_id) %>%
summarize(weekly_driver_revenue = mean(weekly_driver_revenue))

# time driver has been with Lyft
ride_info_drivers <- ride_info_drivers %>%
  group_by(driver_id) %>%
  mutate(driver_onboard_date2 = max(driver_onboard_date2),
         last_ride_date = max(timestamp2)) %>%
  mutate(days_with_lyft = last_ride_date - driver_onboard_date2) %>%
  ungroup()

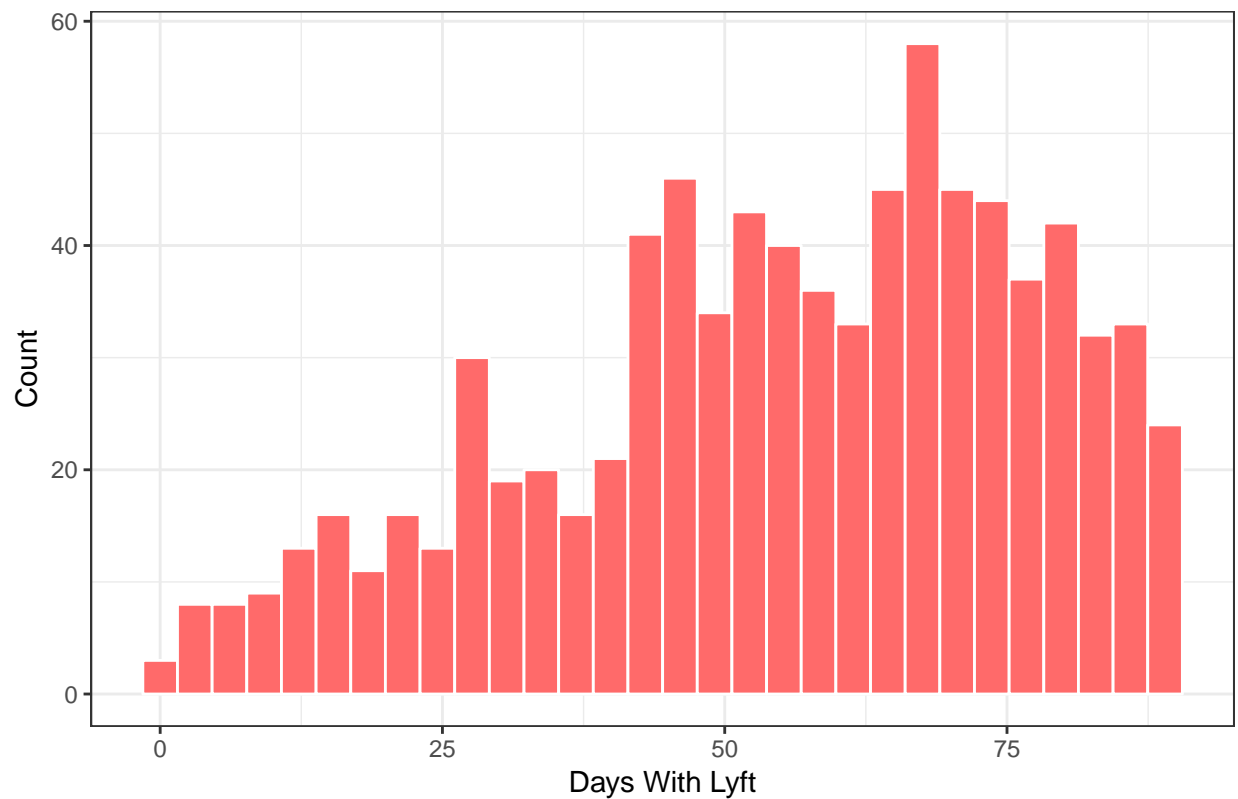
days_with_lyft <- ride_info_drivers %>%
  select(driver_id, days_with_lyft) %>%
  group_by(driver_id) %>%
  summarize(days_with_lyft = max(days_with_lyft)) %>%
  ungroup()

# -----
# distribution of drivers' days with lyft
ride_info_drivers %>%
  group_by(driver_id) %>%
  summarize(days_with_lyft = max(days_with_lyft)) %>%
  mutate(days_with_lyft = as.numeric(days_with_lyft)) %>%
  ungroup() %>%
  ggplot() +
  geom_histogram(aes(x = days_with_lyft),
                 color = "white",
                 fill = "indianred1"
                 ) +
  labs(title = "Distribution of Drivers' Days with Lyft") +
  xlab("Days With Lyft") +
  ylab("Count")

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 18 rows containing non-finite values (stat_bin).

```

Distribution of Drivers' Days with Lyft

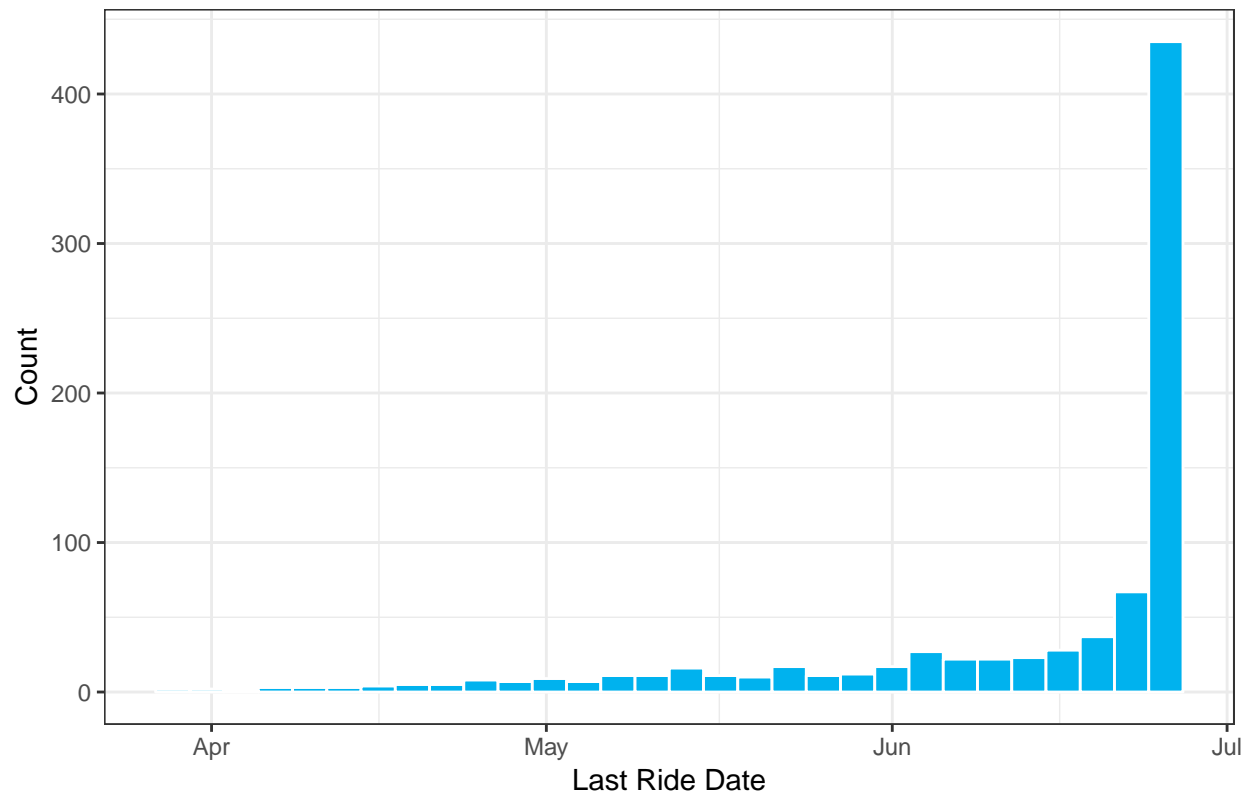


```
ride_info_drivers %>%  
  group_by(driver_id) %>%  
  summarize(last_ride_date = max(last_ride_date)) %>%  
  ggplot(aes(x = last_ride_date)) +  
  geom_histogram(color = "white", fill = "deepskyblue2") +  
  labs(title = "Distribution of Last Ride Dates") +  
  xlab("Last Ride Date") +  
  ylab("Count")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 18 rows containing non-finite values (stat_bin).
```

Distribution of Last Ride Dates




```
# density histogram to visualize retention
date_density <- gather(ride_info_drivers, type, date, driver_onboard_date2, last_ride_date)

date_density <- date_density %>%
  select(driver_id, type, date)

date_density %>%
  group_by(driver_id, type) %>%
  summarize(date = max(date)) %>%
  ggplot(aes(color = type, fill = type)) +
  geom_density(aes(x = date),
               alpha = 0.2) +
  ylab("Density") +
  xlab("Date") +
  scale_color_discrete(name="",
                      labels=c("Driver Onboard Date",
                               "Last Ride Date")) +
  scale_fill_discrete(name="",
                     labels=c("Driver Onboard Date",
                              "Last Ride Date"))

## Warning: Removed 18 rows containing non-finite values (stat_density).
```

```
# -----

# number of rides
num_rides <- ride_info_drivers %>%
  group_by(driver_id) %>%
  distinct(ride_id) %>%
  summarize(num_rides = n())

# mean requested times
requested <- ride_info_drivers %>%
  select(driver_id, event, times) %>%
  filter(event == "requested_at") %>%
  group_by(driver_id) %>%
  summarize(requested_mean = mean(times),
            requested_total = sum(times)) %>%
  ungroup()

# mean accepted times
accepted <- ride_info_drivers %>%
  select(driver_id, event, times) %>%
  filter(event == "accepted_at") %>%
  group_by(driver_id) %>%
  summarize(accepted_mean = mean(times),
            accepted_total = sum(times)) %>%
  ungroup()
```

```

# mean arrived times
arrived <- ride_info_drivers %>%
  select(driver_id, event, times) %>%
  filter(event == "arrived_at") %>%
  group_by(driver_id) %>%
  summarize(arrived_mean = mean(times),
            arrived_total = sum(times)) %>%
  ungroup()

# mean picked up times
picked_up <- ride_info_drivers %>%
  select(driver_id, event, times) %>%
  filter(event == "picked_up_at") %>%
  group_by(driver_id) %>%
  summarize(picked_up_mean = mean(times),
            picked_up_total = sum(times)) %>%
  ungroup()

# mean dropped off times
dropped_off <- ride_info_drivers %>%
  select(driver_id, event, times) %>%
  filter(event == "dropped_off_at") %>%
  group_by(driver_id) %>%
  summarize(dropped_off_mean = mean(times),
            dropped_off_total = sum(times)) %>%
  ungroup()

# requested-arrived time lapse
requested_arrived <- full_join(requested, arrived, by = "driver_id")

requested_arrived <- requested_arrived %>%
  mutate(requested_arrived = arrived_mean - requested_mean) %>%
  select(driver_id, requested_arrived)

# accepted - arrived time lapse

accepted_arrived <- full_join(accepted, arrived, by = "driver_id")

accepted_arrived <- accepted_arrived %>%
  mutate(accepted_arrived = arrived_mean - accepted_mean) %>%
  select(driver_id, accepted_arrived)

# requested - dropped off time lapse
# ride duration is only based on the arrived - dropped off
requested_dropped_off <- full_join(requested, dropped_off, by = "driver_id")

requested_dropped_off <- requested_dropped_off %>%
  mutate(requested_dropped_off = dropped_off_mean - requested_mean) %>%
  select(driver_id, requested_dropped_off)

# day/night
day_night <- ride_info_drivers %>%
  filter(event == "accepted_at") %>%

```

```

select(driver_id, timestamp3) %>%
mutate(hour = hour(timestamp3)) %>%
group_by(driver_id) %>%
summarize(hour = mean(hour)) %>%
ungroup() %>%
mutate(day_night = case_when((hour >= 6 & hour < 12) ~ "morning",
                             (hour >= 12 & hour < 18) ~ "afternoon",
                             (hour < 6 ) ~ "night",
                             TRUE ~ "evening"
                             )) %>%

select(driver_id, day_night)

# average miles per hour
mph <- ride_info_drivers %>%
mutate(ride_duration_hour = ride_duration_min / 60) %>%
group_by(driver_id) %>%
mutate(mean_mph = ride_distance_miles / ride_duration_hour) %>%
summarize(mean_mph = mean(mean_mph)) %>%
ungroup()

# counts of prime time
prime_time_counts <- ride_info_drivers %>%
filter(ride_prime_time != 0) %>%
group_by(driver_id) %>%
distinct(ride_id) %>%
summarize(prime_time_counts = n()) %>%
ungroup()

# Graphing total number of rides vs. the number of rides
# that had Prime Time applied to them
num_rides_vs_num_prime <- full_join(num_rides,
                                     prime_time_counts,
                                     by = "driver_id")

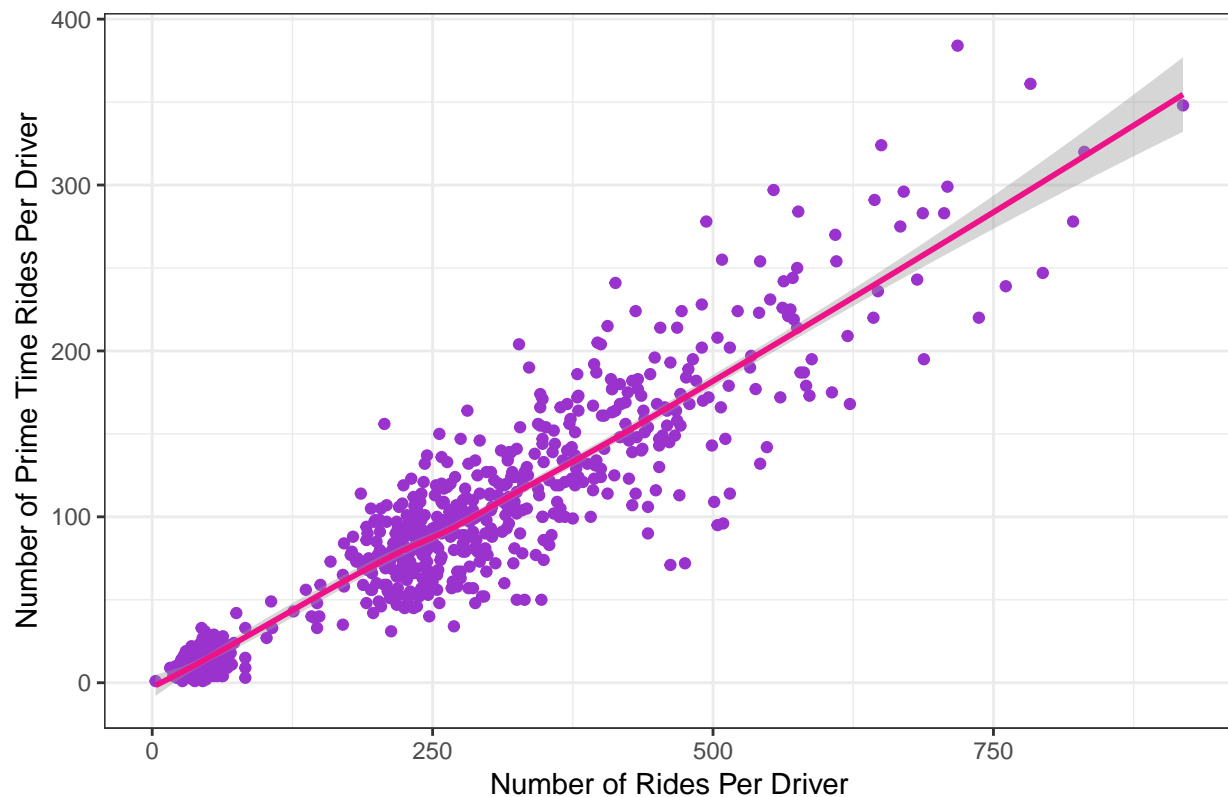
num_rides_vs_num_prime %>%
  ggplot() +
  geom_point(aes(x = num_rides, y = prime_time_counts),
             color = "darkorchid3") +
  geom_smooth(aes(x = num_rides, y = prime_time_counts),
             method = "loess", color = "deeppink2") +
  labs(title = "Total of Rides Vs. Total of Prime Time Rides") +
  xlab("Number of Rides Per Driver") +
  ylab("Number of Prime Time Rides Per Driver")

```

```
## Warning: Removed 3 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 3 rows containing missing values (geom_point).
```

Total of Rides Vs. Total of Prime Time Rides



```
# -----
# driver revenue dataset
driver_revenue <- ride_info_drivers %>%
  group_by(driver_id) %>%
  summarize(driver_revenue = max(driver_revenue))

driver_revenue <- left_join(driver_revenue, days_with_lyft)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, num_rides)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, requested)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, accepted)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, arrived)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, picked_up)

## Joining, by = "driver_id"
```

```

driver_revenue <- left_join(driver_revenue, dropped_off)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, requested_arrived)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue,
                             accepted_arrived)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue,
                             requested_dropped_off)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, day_night)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue, mph)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue,
                             prime_time_counts)

## Joining, by = "driver_id"
driver_revenue <- left_join(driver_revenue,
                             weekly_driver_revenue)

## Joining, by = "driver_id"
driver_revenue <- driver_revenue %>%
  mutate(num_prime_num_rides = prime_time_counts / num_rides)

### Statistical Tests (Parametric & Non-Parametric)

# simple correlation tests in order to see what factors
# matter the most when it comes to driver revenue
# r - linear correlation (+1 is a strong positive relationship,
#                          -1 is a strong negative relationship,
#                          0 means no relationship)
# R^2 - how much of the variation of the dependent variable
#       can be attributed to the independent variable

## Factors Being Considered
# days_with_lyft
# num_rides
# requested
# accepted
# arrived
# picked_up
# dropped_off
# requested_arrived

```

```

# accepted_arrived
# requested_dropped_off
## day_night <- not being included with correlation
# mean_mph
# prime_time_counts
# num_prime_num_rides

# removing NAs
driver_revenue2 <- na.omit(driver_revenue)

# calculating the r-sq values in order to represent how much the
# variation in driver revenue can be explained
# by each individual factor
rsq <- rep(0,13)
p_values <- rep(0,13)

possible_features <- c("Days With Lyft",
                      "Number of Rides",
                      "Average Requested Time",
                      "Average Accepted Time",
                      "Average Arrived Time",
                      "Average Picked Up Time",
                      "Average Dropped Off Time",
                      "Average Requested-Arrived Timelapse",
                      "Average Accepted-Arrived Timelapse",
                      "Average Requested-Dropped Off Timelapse",
                      "Average Speed (MPH)",
                      "Number of Prime Time Rides",
                      "Prime Time Rides:Normal Rides Ratio")

rsq[1] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$days_with_lyft)))^2
rsq[2] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$num_rides)))^2
rsq[3] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_mean)))^2
rsq[4] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$accepted_mean)))^2
rsq[5] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$arrived_mean)))^2
rsq[6] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$picked_up_mean)))^2
rsq[7] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$dropped_off_mean)))^2
rsq[8] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_arrived)))^2
rsq[9] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$accepted_arrived)))^2
rsq[10] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_dropped_off)))^2
rsq[11] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$mean_mph)))^2
rsq[12] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$prime_time_counts)))^2
rsq[13] <- (cor(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$num_prime_num_rides)))^2

# p-values of each correlation test
# all of them are statistically significant
p_values[1] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$days_with_lyft))$p.value
p_values[2] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$num_rides))$p.value
p_values[3] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_mean))$p.value
p_values[4] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$accepted_mean))$p.value
p_values[5] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$arrived_mean))$p.value
p_values[6] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$picked_up_mean))$p.value

```

```

p_values[7] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$dropped_off_mean))$p
p_values[8] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_arrived))$p
p_values[9] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$accepted_arrived))$p
p_values[10] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$requested_dropped_off))$p
p_values[11] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$mean_mph))$p.value
p_values[12] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$prime_time_counts))$p.value
p_values[13] <- cor.test(driver_revenue2$driver_revenue, as.numeric(driver_revenue2$num_prime_num_rides))$p.value

cor_dfr <- data_frame(possible_features, rsq, p_values)

```

```


## Warning: `data_frame()` is deprecated, use `tibble()`.
## This warning is displayed once per session.

```

```

# bar plot showing each factors Rsq value
cor_dfr %>%
  ggplot() +
  geom_col(aes(x = reorder(possible_features, rsq), y = rsq, fill = possible_features)) +
  coord_flip() +
  ylab("R-Squared Values") +
  xlab("Possible Features") +
  labs(title = "Two-Sided Correlation Tests") +
  theme(legend.position = "none")

```



```

# plot visualizing the significant correlation between the number
# of rides have on the variation in driver revenue
driver_revenue %>%
  mutate(num_rides_group = case_when((num_rides < 200) ~ "0 - 200",

```

```


      (num_rides < 400) ~ "200 - 400",
      (num_rides < 600) ~ "400 - 600",
      (num_rides < 800) ~ "600 - 800",
      TRUE ~ "800+"
    )
  ) %>%
  mutate(num_rides_group = as.factor(num_rides_group)) %>%
  ggplot(aes(x = days_with_lyft, y = driver_revenue)) +
  geom_point(aes(color = num_rides_group)) +
  ylab("Driver Revenue, Dollars") +
  xlab("Days with Lyft") +
  geom_smooth(aes(x = days_with_lyft, y = driver_revenue), group = 1, alpha = 0.25) +
  scale_y_continuous(labels = comma) +
  labs(color = "Number of Rides")

```

```

## Don't know how to automatically pick scale for object of type difftime. Defaulting to continuous.
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
## Warning: Removed 18 rows containing non-finite values (stat_smooth).
## Warning: Removed 18 rows containing missing values (geom_point).

```



```

# another plot visualizing the same thing in another way
driver_revenue %>%
  mutate(days_with_lyft = as.numeric(days_with_lyft)) %>%
  mutate(days_with_lyft_group = case_when(
    (days_with_lyft < 20) ~ "0 - 20",
    (days_with_lyft < 40) ~ "20 - 40",
    (days_with_lyft < 60) ~ "40 - 60",

```



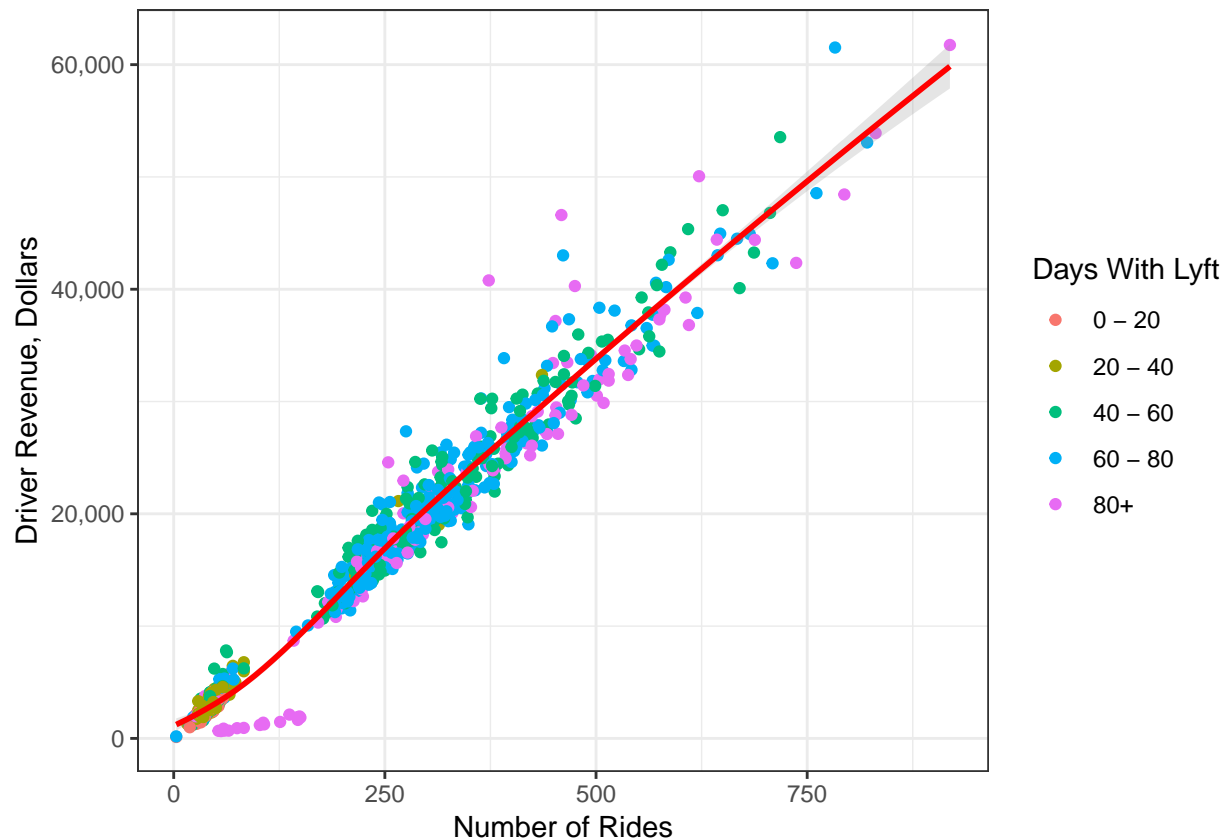
```

(days_with_lyft < 80) ~ "60 - 80",
TRUE ~ "80+"

)) %>%
mutate(days_with_lyft_group = as.factor(days_with_lyft_group)) %>%
ggplot(aes(x = num_rides, y = driver_revenue)) +
geom_point(aes(color = days_with_lyft_group)) +
ylab("Driver Revenue, Dollars") +
xlab("Number of Rides") +
geom_smooth(aes(x = num_rides, y = driver_revenue), group = 1, alpha = 0.25, color = "red") +
scale_y_continuous(labels = comma) +
labs(color = "Days With Lyft")

```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```

revenue <- driver_revenue$driver_revenue
day_night <- as.factor(driver_revenue$day_night)


# visualizing the distributions of driver revenue among
# different groups of drivers
# drivers are grouped by the time of day they are most active
driver_revenue2 %>%
  mutate(day_night = case_when((day_night == "morning") ~ "Morning",
                                (day_night == "afternoon") ~ "Afternoon",
                                (day_night == "evening") ~ "Evening",
                                (day_night == "night") ~ "Night"

```

```

)) %>%
ggplot() +
geom_boxplot(aes(x = as.factor(day_night), y = driver_revenue,
                 fill = as.factor(day_night))) +
theme(axis.text.x = element_text(angle = 45)) +
xlab("Most Active Time of Day") +
ylab("Driver Revenue") +
scale_fill_discrete(name = "Time of Day")

```

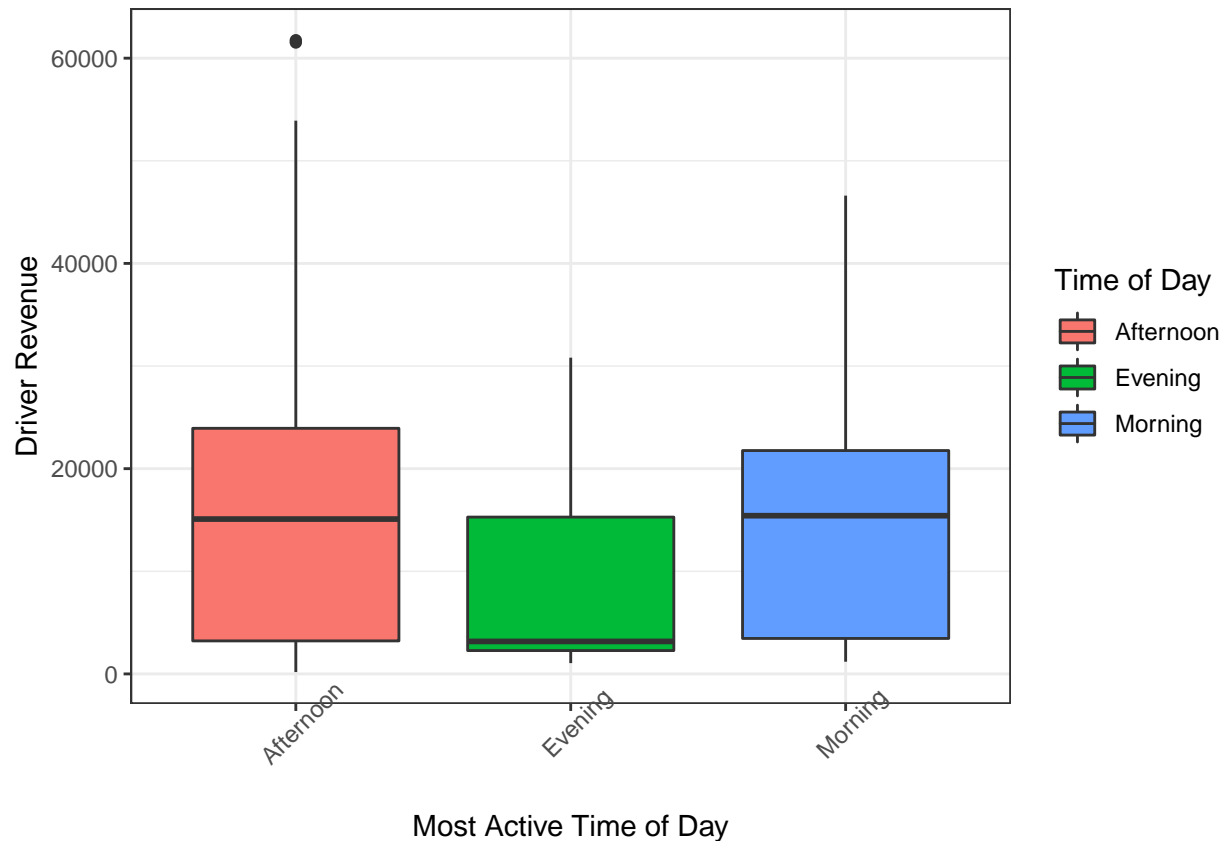


```

# removes 'night' for a cleaner graph
driver_revenue2 %>%
  filter(day_night != "night") %>%
  mutate(day_night = case_when((day_night == "morning") ~ "Morning",
                              (day_night == "afternoon") ~ "Afternoon",
                              (day_night == "evening") ~ "Evening",
                              (day_night == "night") ~ "Night"))

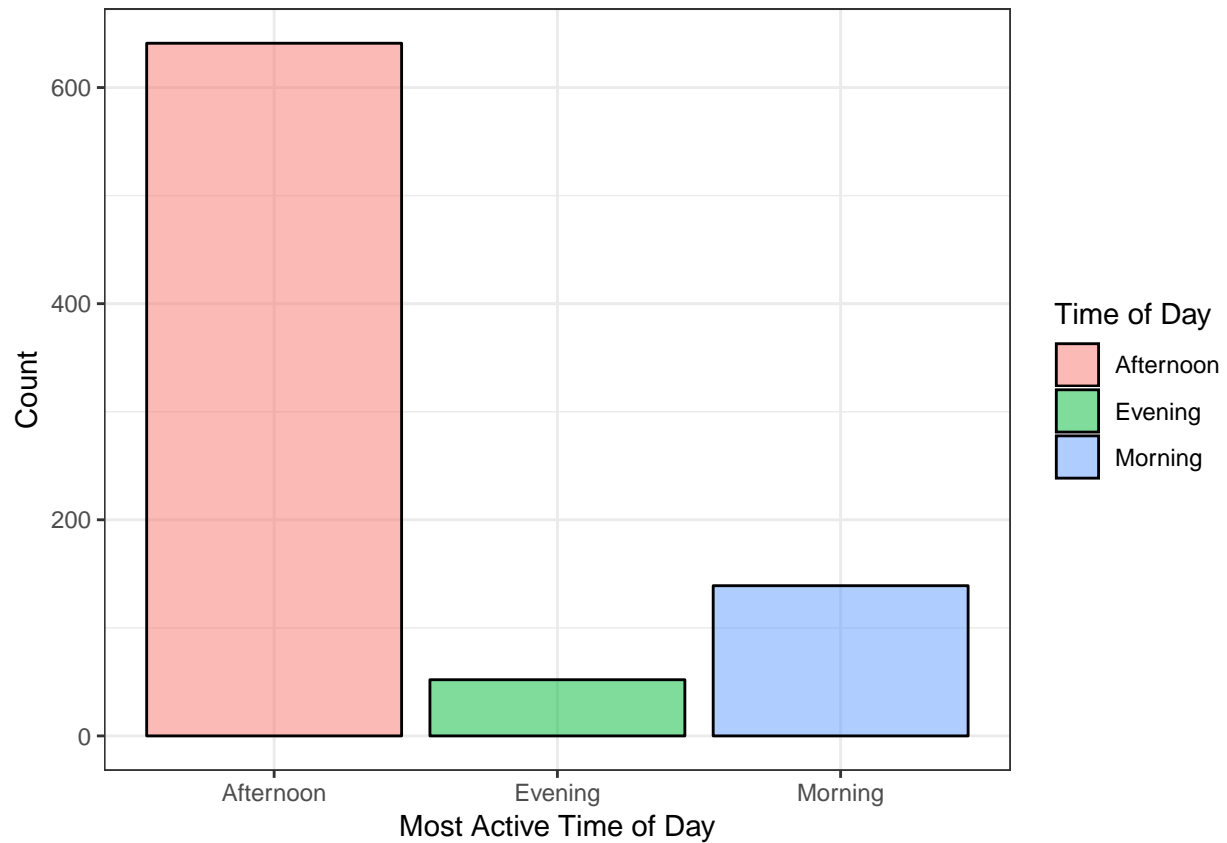
)) %>%
ggplot() +
geom_boxplot(aes(x = as.factor(day_night), y = driver_revenue,
                 fill = as.factor(day_night))) +
theme(axis.text.x = element_text(angle = 45)) +
xlab("Most Active Time of Day") +
ylab("Driver Revenue") +
scale_fill_discrete(name = "Time of Day")

```



```
# two graphs showing how many rides are completed in each time
# of day
driver_revenue2 %>%
  filter(day_night != "night") %>%
  mutate(day_night = case_when((day_night == "morning") ~ "Morning",
                                (day_night == "afternoon") ~ "Afternoon",
                                (day_night == "evening") ~ "Evening",
                                (day_night == "night") ~ "Night"

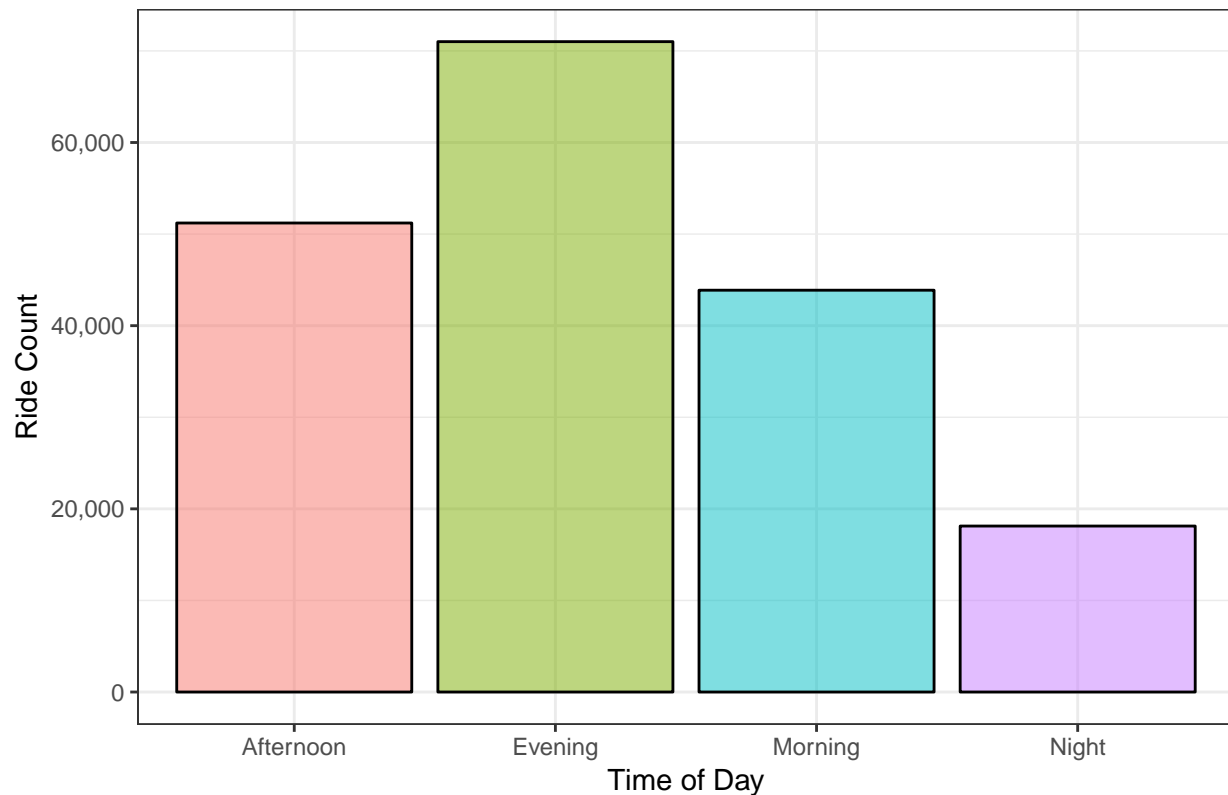
  )) %>%
  ggplot() +
  geom_bar(aes(x = day_night, fill = day_night),
           color = "black",
           alpha = 0.5) +
  xlab("Most Active Time of Day") +
  ylab("Count") +
  scale_fill_discrete(name = "Time of Day")
```



```
ride_info_drivers %>%
  filter(event == "accepted_at") %>%
  select(driver_id, timestamp3) %>%
  mutate(hour = hour(timestamp3)) %>%
  mutate(day_night = case_when((hour >= 6 & hour < 12) ~ "Morning",
                                (hour >= 12 & hour < 18) ~ "Afternoon",
                                (hour < 6 ) ~ "Night",
                                TRUE ~ "Evening"
                                )) %>%

  ggplot(aes(x = day_night)) +
  geom_bar(aes(fill = day_night),
           color = "black",
           alpha = 0.5) +
  scale_y_continuous(labels = comma) +
  xlab("Time of Day") +
  ylab("Ride Count") +
  theme(legend.position = "none") +
  ggtitle("Number of Rides By Time of Day")
```

Number of Rides By Time of Day



```
# compact dataframe is made to create
# a correlation matrix heatmap, to check for interdependent
# variables
driver_revenue3 <- driver_revenue2 %>%
  select(num_rides, prime_time_counts, days_with_lyft, day_night) %>%
  mutate(days_with_lyft = as.numeric(days_with_lyft)) %>%
  mutate(day_night = case_when((day_night == "morning") ~ 1,
                                (day_night == "afternoon") ~ 2,
                                (day_night == "evening") ~ 3,
                                (day_night == "night") ~ 4

  ))

cormat <- round(cor(driver_revenue3), 2)

get_lower_tri <- function(cormat){
  cormat[upper.tri(cormat)] <- NA
  return(cormat)
}

# upper triangle of the correlation matrix
get_upper_tri <- function(cormat){
  cormat[lower.tri(cormat)] <- NA
  return(cormat)
}
```

```
upper_tri <- get_upper_tri(cormat)
upper_tri
```


```
##               num_rides prime_time_counts days_with_lyft day_night
## num_rides           1              0.94          0.53      -0.05
## prime_time_counts    NA              1.00          0.45       0.00
## days_with_lyft       NA              NA          1.00      -0.07
## day_night            NA              NA          NA       1.00
```

```
# melt the correlation matrix
```

```
melted_cormat <- melt(upper_tri, na.rm = TRUE)
```

```
# correlation matrix heatmap
```

```
ggplot(data = melted_cormat, aes(Var2, Var1, fill = value))+
  geom_tile(color = "white")+
  scale_fill_gradient2(low = "blue", high = "red", mid = "white",
                      midpoint = 0, limit = c(-1,1), space = "Lab",
                      name="Pearson\nCorrelation") +
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 45, vjust = 1,
                                    size = 12, hjust = 1)) +
  coord_fixed() +
  xlab("") +
  ylab("")
```



```
# linear regression model to calculate projected lifetime of a driver
```

```

# response variable: days_with_lyft

# possible features:
# driver revenue
# average number of rides per week *
# average trip_duration *
# average trip_distance *
# average requested_arrived timelapse
# average accepted_arrived timelapse
# average accepted_requested_dropped_off timelapse
# day_night * (might need to convert to numbers)
# num_prime_rides:num_rides ratio

driver_revenue2 <- driver_revenue2 %>%
  mutate(days_with_lyft = as.numeric(days_with_lyft)) %>%
  mutate(weeks_with_lyft = days_with_lyft / 7) %>%
  mutate(mean_rides_per_week = num_rides / weeks_with_lyft)

total_trip_stats <- ride_info %>%
  group_by(driver_id) %>%
  summarize(total_duration = sum(ride_duration),
            total_distance = sum(ride_distance))

driver_revenue2 <- full_join(total_trip_stats, driver_revenue2,
                             by = "driver_id")

driver_revenue2 <- driver_revenue2 %>%
  mutate(mean_duration = total_duration / num_rides,
         mean_distance = total_distance / num_rides)

model1 <- lm(days_with_lyft ~ (weekly_driver_revenue +
                              mean_rides_per_week +
                              (mean_duration *
                               mean_distance *
                               requested_dropped_off) +
                              (requested_arrived *
                               accepted_arrived) +
                              day_night +
                              num_prime_num_rides
                              ),
  data = driver_revenue2)

summary(model1)

##
## Call:
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +
##   (mean_duration * mean_distance * requested_dropped_off) +
##   (requested_arrived * accepted_arrived) + day_night + num_prime_num_rides),
##   data = driver_revenue2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.397 -14.059   1.257  14.852  63.974

```

```
##
## Coefficients:
##
##               Estimate Std. Error
## (Intercept)    3.805e+01  1.742e+01
## weekly_driver_revenue    5.412e-02  4.682e-03
## mean_rides_per_week   -5.498e-01  5.138e-02
## mean_duration         6.081e-03  4.063e-03
## mean_distance       -2.913e-05  4.384e-04
## requested_dropped_off    1.746e-02  8.150e-03
## requested_arrived   -1.704e-02  1.765e-02
## accepted_arrived     9.505e-03  1.794e-02
## day_nightevening     7.642e-01  2.915e+00
## day_nightmorning     3.869e+00  1.867e+00
## day_nightnight    -1.470e+01  1.952e+01
## num_prime_num_rides  -1.103e+01  7.055e+00
## mean_duration:mean_distance  -6.171e-08  9.191e-08
## mean_duration:requested_dropped_off  -4.817e-06  1.994e-06
## mean_distance:requested_dropped_off  -1.899e-07  2.168e-07
## requested_arrived:accepted_arrived  -4.802e-06  1.271e-06
## mean_duration:mean_distance:requested_dropped_off  5.762e-11  4.445e-11
##
##               t value Pr(>|t|)
## (Intercept)      2.184 0.029264 *
## weekly_driver_revenue    11.560 < 2e-16 ***
## mean_rides_per_week   -10.700 < 2e-16 ***
## mean_duration         1.496 0.134912
## mean_distance       -0.066 0.947047
## requested_dropped_off    2.143 0.032428 *
## requested_arrived   -0.965 0.334604
## accepted_arrived     0.530 0.596471
## day_nightevening     0.262 0.793273
## day_nightmorning     2.072 0.038593 *
## day_nightnight    -0.753 0.451777
## num_prime_num_rides  -1.563 0.118441
## mean_duration:mean_distance  -0.671 0.502178
## mean_duration:requested_dropped_off  -2.416 0.015898 *
## mean_distance:requested_dropped_off  -0.876 0.381305
## requested_arrived:accepted_arrived  -3.779 0.000169 ***
## mean_duration:mean_distance:requested_dropped_off  1.296 0.195225
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.42 on 816 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared:  0.2154, Adjusted R-squared:  0.2001
## F-statistic: 14.01 on 16 and 816 DF, p-value: < 2.2e-16

# highest non-interaction insignificant p-value is
# mean_distance, removing it for model 2

model2 <- lm(days_with_lyft ~ (weekly_driver_revenue +
                             mean_rides_per_week +
                             (mean_duration *
                              requested_dropped_off) +
                             (requested_arrived *
```



```

        accepted_arrived) +
        day_night +
        +
        num_prime_num_rides
),
data = driver_revenue2)

summary(model2)

##
## Call:
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +
##   (mean_duration * requested_dropped_off) + (requested_arrived *
##     accepted_arrived) + day_night + +num_prime_num_rides), data = driver_revenue2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45.731 -14.730   0.723  15.149  59.409
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      5.022e+01  6.583e+00   7.629 6.58e-14
## weekly_driver_revenue      5.255e-02  4.619e-03  11.377 < 2e-16
## mean_rides_per_week     -5.210e-01  5.041e-02 -10.335 < 2e-16
## mean_duration       -1.311e-05  1.461e-03  -0.009  0.9928
## requested_dropped_off      6.172e-03  3.345e-03   1.845  0.0653
## requested_arrived     -1.561e-02  1.768e-02  -0.883  0.3773
## accepted_arrived      7.451e-03  1.797e-02   0.415  0.6784
## day_nightevening     -7.045e-02  2.896e+00  -0.024  0.9806
## day_nightmorning      3.363e+00  1.859e+00   1.809  0.0708
## day_nightnight     -1.578e+01  1.958e+01  -0.806  0.4205
## num_prime_num_rides    -2.206e+00  6.282e+00  -0.351  0.7256
## mean_duration:requested_dropped_off -1.434e-06  8.189e-07  -1.751  0.0803
## requested_arrived:accepted_arrived -5.204e-06  1.267e-06  -4.109 4.37e-05
##
## (Intercept)          ***
## weekly_driver_revenue      ***
## mean_rides_per_week      ***
## mean_duration
## requested_dropped_off      .
## requested_arrived
## accepted_arrived
## day_nightevening
## day_nightmorning          .
## day_nightnight
## num_prime_num_rides
## mean_duration:requested_dropped_off .
## requested_arrived:accepted_arrived ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.49 on 820 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared:  0.206, Adjusted R-squared:  0.1944

```

```
## F-statistic: 17.73 on 12 and 820 DF, p-value: < 2.2e-16
```

```
# highest non-interaction insignificant p-value is  
# mean_duration  
# removing it for model3
```

```
model3 <- lm(days_with_lyft ~ (weekly_driver_revenue +  
                                mean_rides_per_week +  
                                requested_dropped_off +  
                                (requested_arrived *  
                                 accepted_arrived) +  
                                day_night +  
                                num_prime_num_rides  
) ,  
data = driver_revenue2)  
summary(model3)
```

```
##
```

```
## Call:
```

```
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +  
##     requested_dropped_off + (requested_arrived * accepted_arrived) +  
##     day_night + num_prime_num_rides), data = driver_revenue2)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -46.269 -14.387   0.848  14.968  58.810
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)      4.969e+01  2.426e+00  20.481 < 2e-16  
## weekly_driver_revenue      5.264e-02  4.469e-03  11.779 < 2e-16  
## mean_rides_per_week     -5.185e-01  4.915e-02 -10.549 < 2e-16  
## requested_dropped_off      3.529e-04  6.867e-04   0.514  0.6075  
## requested_arrived     -1.763e-02  1.763e-02  -1.000  0.3177  
## accepted_arrived        8.834e-03  1.794e-02   0.493  0.6225  
## day_nightevening        8.687e-02  2.870e+00   0.030  0.9759  
## day_nightmorning        3.233e+00  1.858e+00   1.740  0.0822  
## day_nightnight     -1.483e+01  1.957e+01  -0.758  0.4488  
## num_prime_num_rides     -1.827e+00  6.277e+00  -0.291  0.7711  
## requested_arrived:accepted_arrived -5.600e-06  1.248e-06  -4.486 8.31e-06
```

```
##
```

```
## (Intercept)          ***  
## weekly_driver_revenue      ***  
## mean_rides_per_week      ***  
## requested_dropped_off  
## requested_arrived  
## accepted_arrived  
## day_nightevening  
## day_nightmorning          .  
## day_nightnight  
## num_prime_num_rides  
## requested_arrived:accepted_arrived ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 19.51 on 822 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared: 0.2027, Adjusted R-squared: 0.193
## F-statistic: 20.9 on 10 and 822 DF, p-value: < 2.2e-16

# highest non-interaction insignificant p-value(s) is
# time of day, removing it for model4

model4 <- lm(days_with_lyft ~ (weekly_driver_revenue +
                             mean_rides_per_week +
                             requested_dropped_off +
                             (requested_arrived *
                              accepted_arrived) +
                             num_prime_num_rides
                             ),
data = driver_revenue2)

summary(model4)

##
## Call:
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +
## requested_dropped_off + (requested_arrived * accepted_arrived) +
## num_prime_num_rides), data = driver_revenue2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -47.260 -14.438   0.753  15.130  58.807
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      5.031e+01  2.371e+00  21.214 < 2e-16
## weekly_driver_revenue  5.303e-02  4.427e-03  11.979 < 2e-16
## mean_rides_per_week   -5.221e-01  4.912e-02 -10.629 < 2e-16
## requested_dropped_off  4.760e-04  6.825e-04   0.697  0.486
## requested_arrived    -1.732e-02  1.764e-02  -0.982  0.326
## accepted_arrived      8.694e-03  1.794e-02   0.485  0.628
## num_prime_num_rides   -2.547e+00  6.214e+00  -0.410  0.682
## requested_arrived:accepted_arrived -5.423e-06  1.245e-06  -4.355  1.5e-05
##
## (Intercept)          ***
## weekly_driver_revenue ***
## mean_rides_per_week  ***
## requested_dropped_off
## requested_arrived
## accepted_arrived
## num_prime_num_rides
## requested_arrived:accepted_arrived ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.51 on 825 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared: 0.1991, Adjusted R-squared: 0.1923
```

```
## F-statistic: 29.3 on 7 and 825 DF, p-value: < 2.2e-16
# highest non-interaction insignificant p-value is
# the ratio of the number of prime time rides to number of
# non-prime time rides , removing it for model5

model5 <- lm(days_with_lyft ~ (weekly_driver_revenue +
                               mean_rides_per_week +
                               requested_dropped_off +
                               (requested_arrived *
                                accepted_arrived)
                               ),
             data = driver_revenue2)

summary(model5)

##
## Call:
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +
##    requested_dropped_off + (requested_arrived * accepted_arrived)),
##    data = driver_revenue2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.745 -14.525   0.787  15.065  59.063
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      4.957e+01  1.532e+00  32.357 < 2e-16
## weekly_driver_revenue  5.268e-02  4.343e-03  12.130 < 2e-16
## mean_rides_per_week   -5.212e-01  4.904e-02 -10.627 < 2e-16
## requested_dropped_off  4.866e-04  6.816e-04   0.714  0.475
## requested_arrived    -1.739e-02  1.763e-02  -0.987  0.324
## accepted_arrived      8.890e-03  1.792e-02   0.496  0.620
## requested_arrived:accepted_arrived -5.359e-06  1.235e-06  -4.339 1.61e-05
##
## (Intercept)          ***
## weekly_driver_revenue ***
## mean_rides_per_week  ***
## requested_dropped_off
## requested_arrived
## accepted_arrived
## requested_arrived:accepted_arrived ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.5 on 826 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared:  0.199, Adjusted R-squared:  0.1931
## F-statistic: 34.19 on 6 and 826 DF, p-value: < 2.2e-16
# highest non-interaction insignificant p-value is
# the requested-dropped off timelapse, removing it for model6

model6 <- lm(days_with_lyft ~ (weekly_driver_revenue +
```

```

        mean_rides_per_week +
        (requested_arrived *
         accepted_arrived)
),
data = driver_revenue2)

summary(model6)

##
## Call:
## lm(formula = days_with_lyft ~ (weekly_driver_revenue + mean_rides_per_week +
##   (requested_arrived * accepted_arrived)), data = driver_revenue2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.395 -14.839   0.681  14.919  59.295
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      4.950e+01  1.529e+00  32.383 < 2e-16
## weekly_driver_revenue  5.276e-02  4.341e-03  12.154 < 2e-16
## mean_rides_per_week   -5.209e-01  4.903e-02 -10.626 < 2e-16
## requested_arrived     -1.601e-02  1.752e-02  -0.914  0.361
## accepted_arrived       8.609e-03  1.791e-02   0.481  0.631
## requested_arrived:accepted_arrived -5.148e-06  1.199e-06  -4.295 1.96e-05
##
## (Intercept)          ***
## weekly_driver_revenue      ***
## mean_rides_per_week       ***
## requested_arrived
## accepted_arrived
## requested_arrived:accepted_arrived ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.5 on 827 degrees of freedom
## (104 observations deleted due to missingness)
## Multiple R-squared:  0.1985, Adjusted R-squared:  0.1936
## F-statistic: 40.96 on 5 and 827 DF,  p-value: < 2.2e-16

# despite there being non-significant p-values present for the
# requested-arrived and accepted-arrived timelapses,
# their interaction remains highly statistically significant

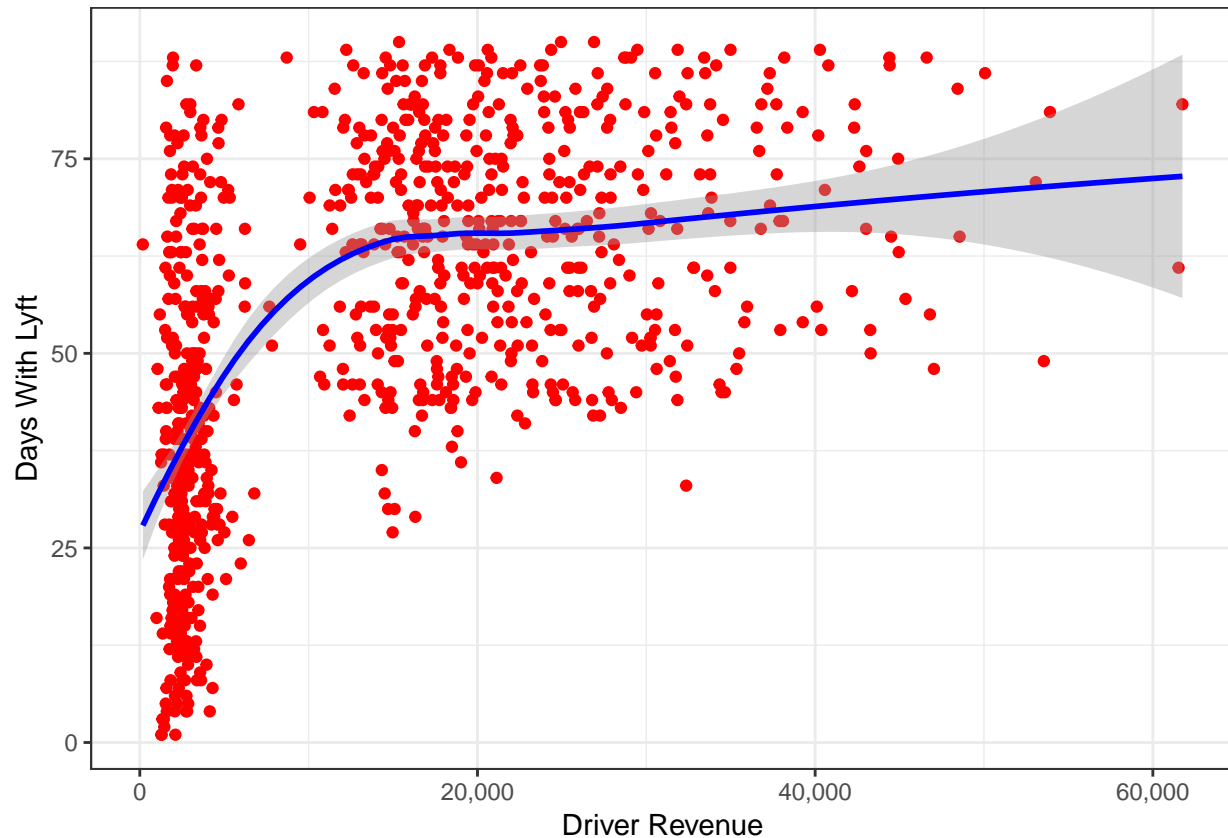
# weekly driver_revenue has the largest coefficient in the
# model, making it the biggest factor on how many days a
# driver stays with lyft
# # graphing weekly_driver_revenue against days_with_lyft
# # to see if polynomial regression is necessary.
# # a loess fit line is used as a guide
#
driver_revenue2 %>%
  arrange(days_with_lyft) %>%
  ggplot(aes(x = driver_revenue, y = days_with_lyft)) +

```

```
geom_point(color = "red") +
geom_smooth(color = "blue",
            method = "loess") +
xlab("Driver Revenue") +
ylab("Days With Lyft") +
scale_x_continuous(labels = comma)
```

Warning: Removed 104 rows containing non-finite values (stat_smooth).

Warning: Removed 104 rows containing missing values (geom_point).



```
#
# # since the loess fit line has heavy curvature polynomial
# # regression will allow the model to represent the data
# # better
#
#
#
```

```
# in order to determine which order to use for polynomial
# regression, k-fold validation is performed.
# For five loops the data is randomly split into 5 sections.
# In each loop, four sections is used to train the model
# and one section is used to validate the model
# the RMSE value is taken to assess model performance
# and for each polynomial order, the average RMSE is recorded
```

```

# the model with the lowest RMSE is recommended

sampled_data <- sample_frac(driver_revenue2,1) %>%
  mutate(count = row_number())

fold <- 5
k <- 10
average_rmse <- rep(0,k)
average_se <- rep(0,k)

for(i in 1:k){

  rmse <- rep(0,fold)
  for(j in 1:fold){

    training_setj <- sampled_data %>%
      mutate(cross = (row_number() + 5) %% 5) %>%
      filter(cross != j - 1)

    validation_setj <- sampled_data %>%
      mutate(cross = (row_number() + 5) %% 5) %>%
      filter(cross == j - 1)

    rmse[j] <- rmse(lm(days_with_lyft ~ (poly(weekly_driver_revenue, i, raw = TRUE) +
                                             mean_rides_per_week +
                                             (requested_arrived *
                                              accepted_arrived)
                                             ),
    data = training_setj), validation_setj)

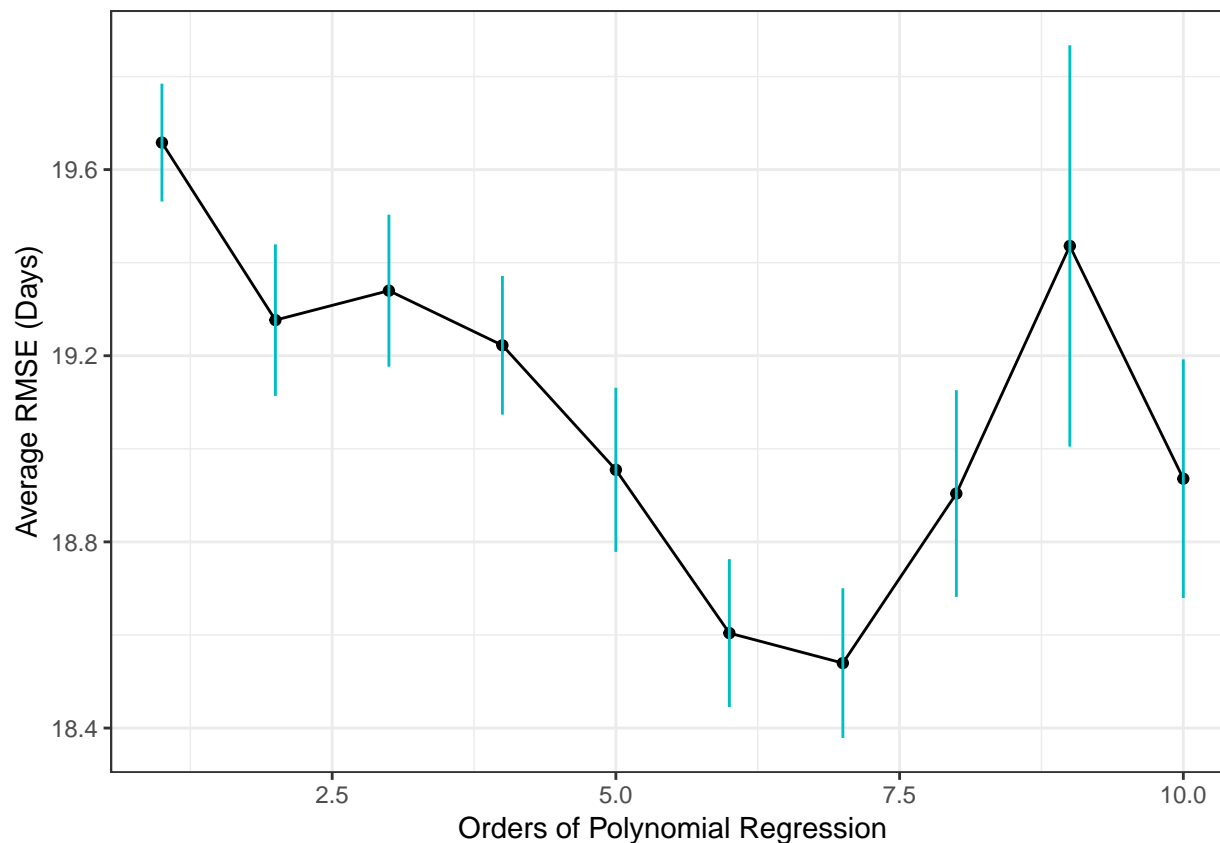
  }

  average_rmse[i] <- mean(rmse)
  average_se[i] <- sd(rmse)/length(rmse)
}

orders <- c(1:10)

data.frame(orders, average_rmse, average_se) %>%
  ggplot(aes(x = orders)) +
  geom_point(aes(y = average_rmse), color = 'black') +
  geom_line(aes(y = average_rmse), color = 'black') +
  geom_linerange(aes(ymin = average_rmse - average_se,
                    ymax = average_rmse + average_se), color = '#00BFC4') +
  ylab("Average RMSE (Days)") +
  xlab("Orders of Polynomial Regression")

```



```
match(min(average_rmse), average_rmse)
```

```
## [1] 7
```

```
average_rmse[match(min(average_rmse), average_rmse)]
```

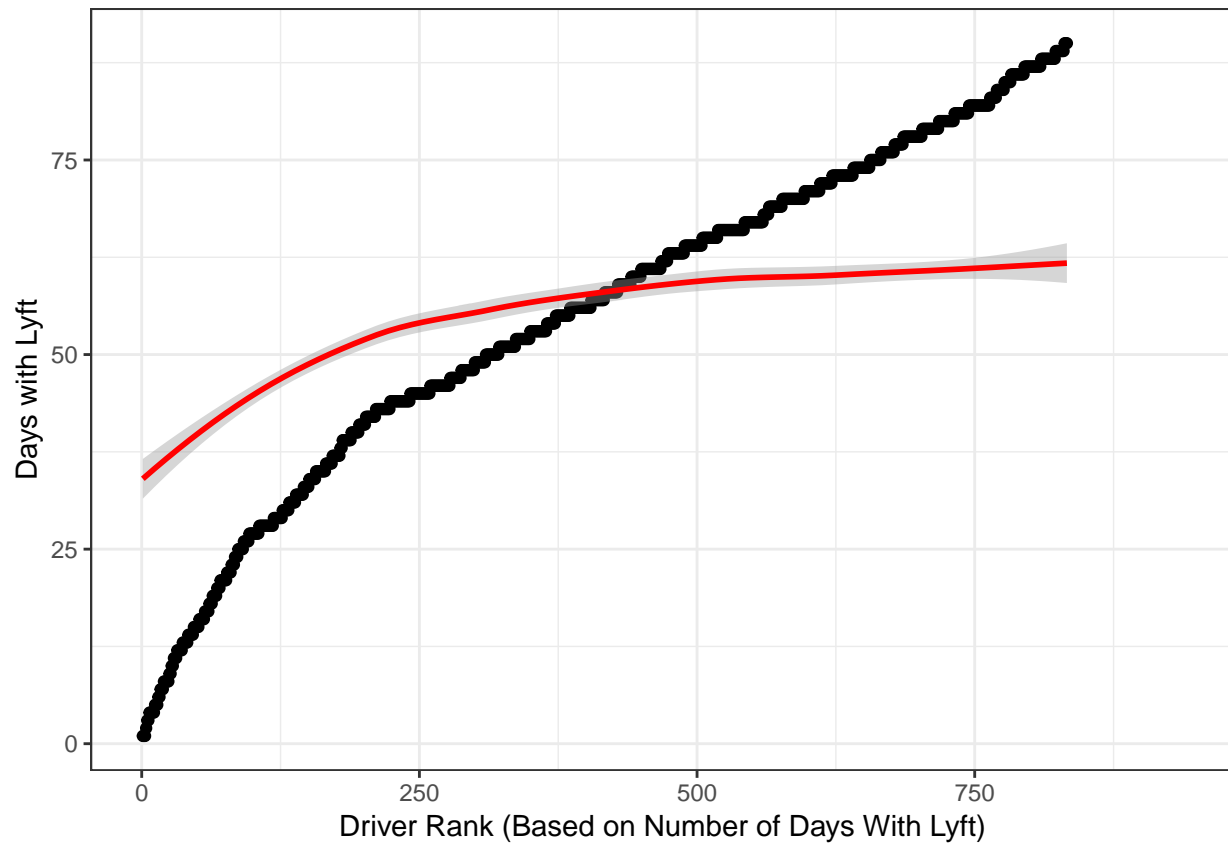
```
## [1] 18.53955
```

```
driver_revenue2 %>%
  arrange(days_with_lyft) %>%
  mutate(rank = row_number()) %>%
  add_predictions(lm(days_with_lyft ~ (poly(weekly_driver_revenue, 7, raw = TRUE) +
    mean_rides_per_week +
    (requested_arrived *
    accepted_arrived)
  ),
  data = driver_revenue2)) %>%
  ggplot(aes(x = rank, y = days_with_lyft)) +
  geom_point() +
  geom_smooth(aes(y = pred), color = "red") +
  ylab("Days with Lyft") +
  xlab("Driver Rank (Based on Number of Days With Lyft)")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 104 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 104 rows containing missing values (geom_point).
```

```
# In order to calculate average driver's lifetime, we
# applied a total of 180 variable scenarios over the
# timeframe of 55 days, and
# attempted to account for the outcomes of the total amount
# of interactions over a span of 45 years (the expected
# work duration of any individual from the age of 20 - 65)
```

```
driver_revenue2 %>%
  add_predictions(lm(days_with_lyft ~ (poly(weekly_driver_revenue, 7, raw = TRUE) +
    mean_rides_per_week +
    (requested_arrived *
    accepted_arrived)
  ),
  data = driver_revenue2)) %>%
  summarize(projected_lifetime = mean(pred, na.rm = TRUE)) %>%
  mutate(projected_lifetime = projected_lifetime * 180)
```

```
## # A tibble: 1 x 1
##   projected_lifetime
##             <dbl>
## 1             9938.
```

```
# Projected Driver Lifetime: ~9938 days or ~27 years
```